

Implementing Simple Regression Models in R



Vitthal Srinivasan

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Build regression models in R


Understand and test the regression assumptions

Use simple regression models in R

- **to explain variance**
- **to make forecasts**

Avoid some common regression pitfalls

Ease of Prototyping



Excel is the fastest prototyping tool out there

Robustness and Reuse



No free lunches

“Make the common use-case easy
and the difficult use-case possible.”

Regression: Excel, R or Python?



Excel

Create a regression slide for an important presentation



R

Create a regression case study for a seminar



Python

Build trading model that scrapes websites, combines sentiment analysis and regression

Regression: Excel, R or Python?



Excel

Presentations



R

Seminars



Python

Trading models

R for Regression



R

Presentations



R

Seminars



R

Trading models

Use **R for regression**: It makes sense whatever your use-case

Excel or R for Linear Regression?

**Number of independent
variables**

Simple or Multiple Regression?

**Sanity check regression
model**

Accept model as-is or analyse and
tweak?

Excel or R for Linear Regression?

**Sanity Check
Result**

Regression Type

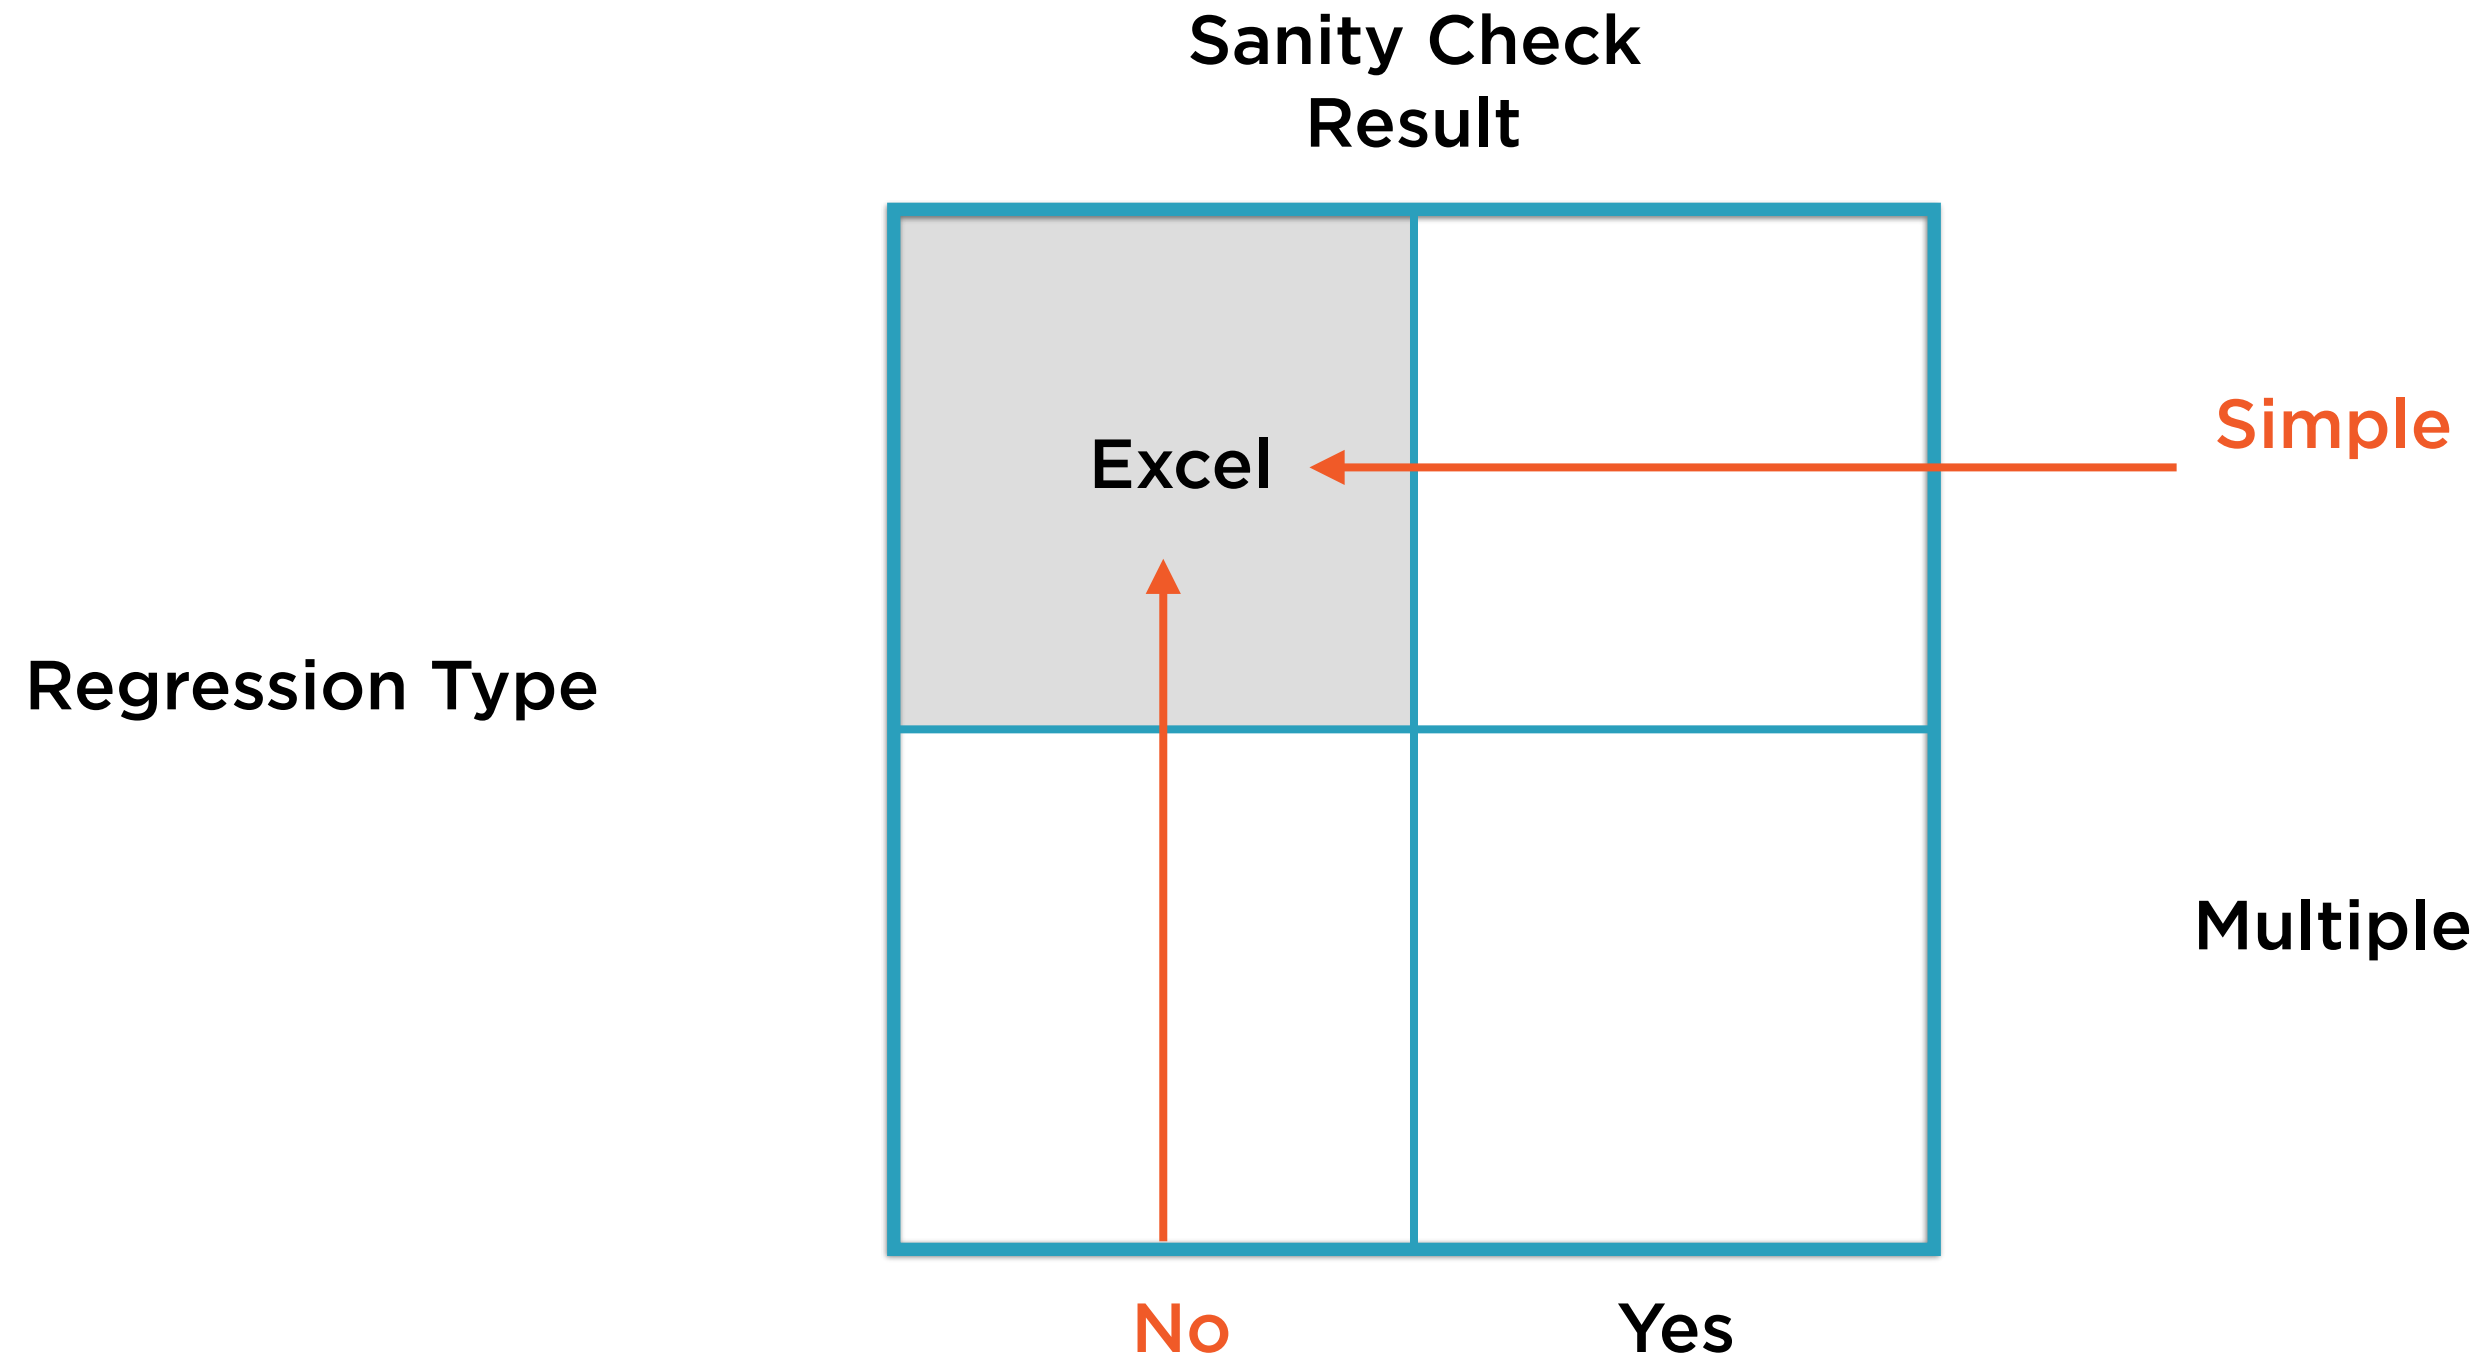
Simple

Multiple

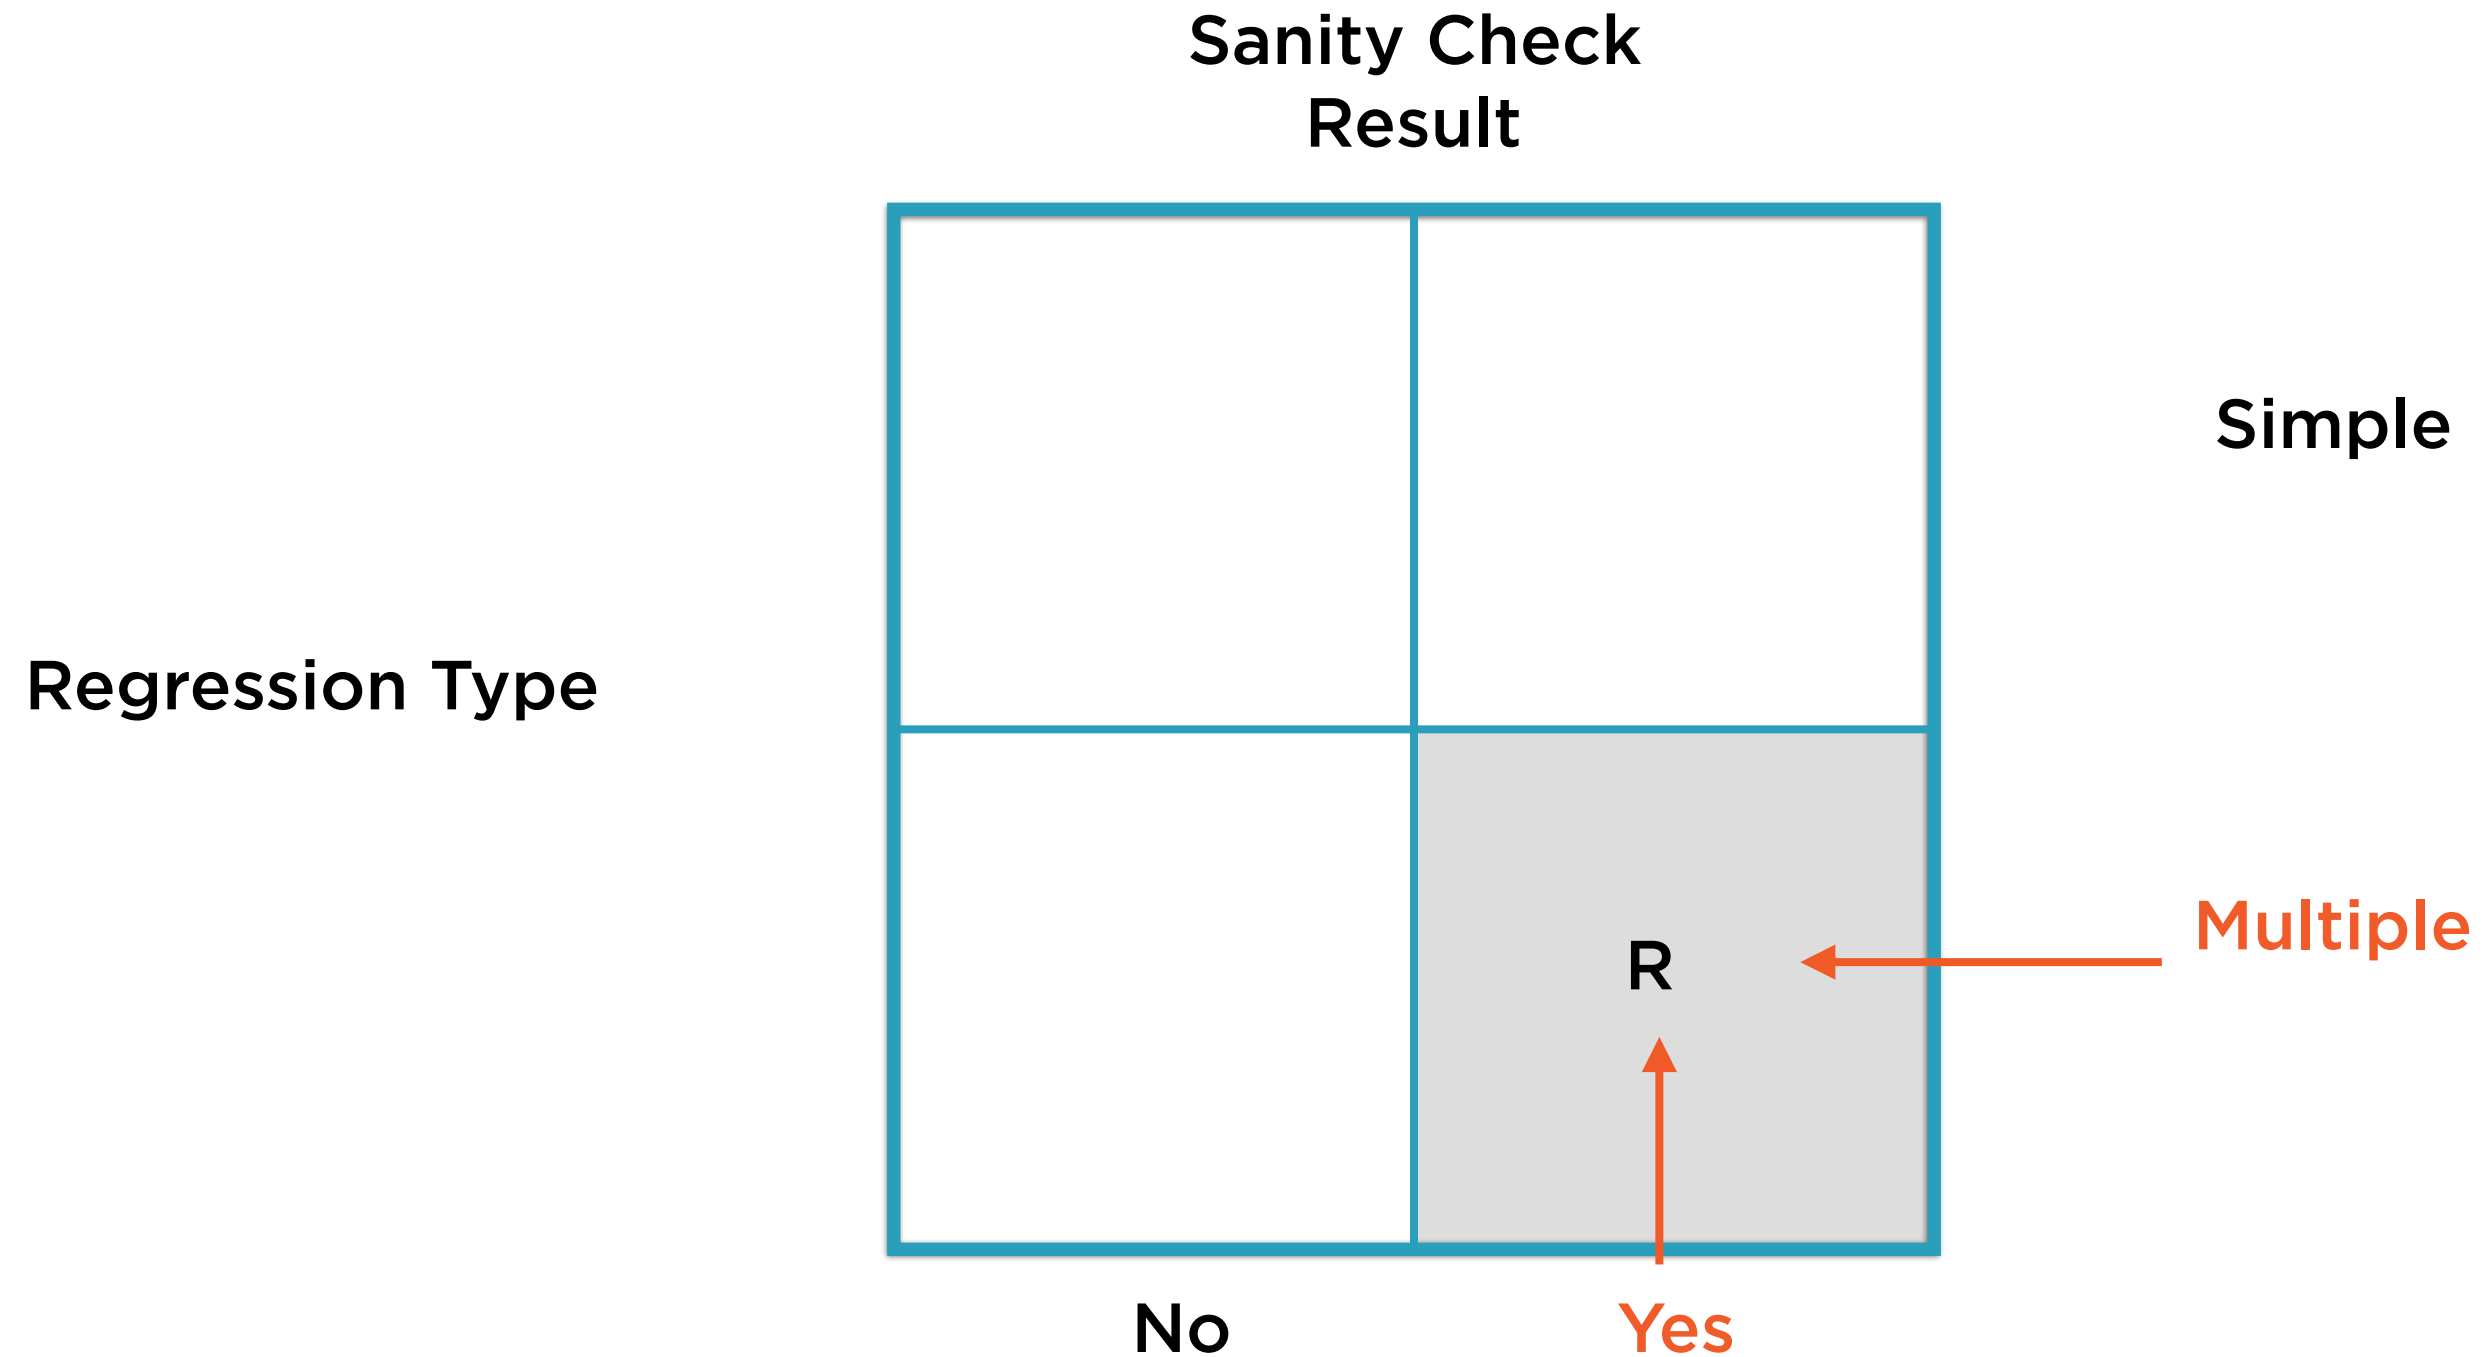
No

Yes

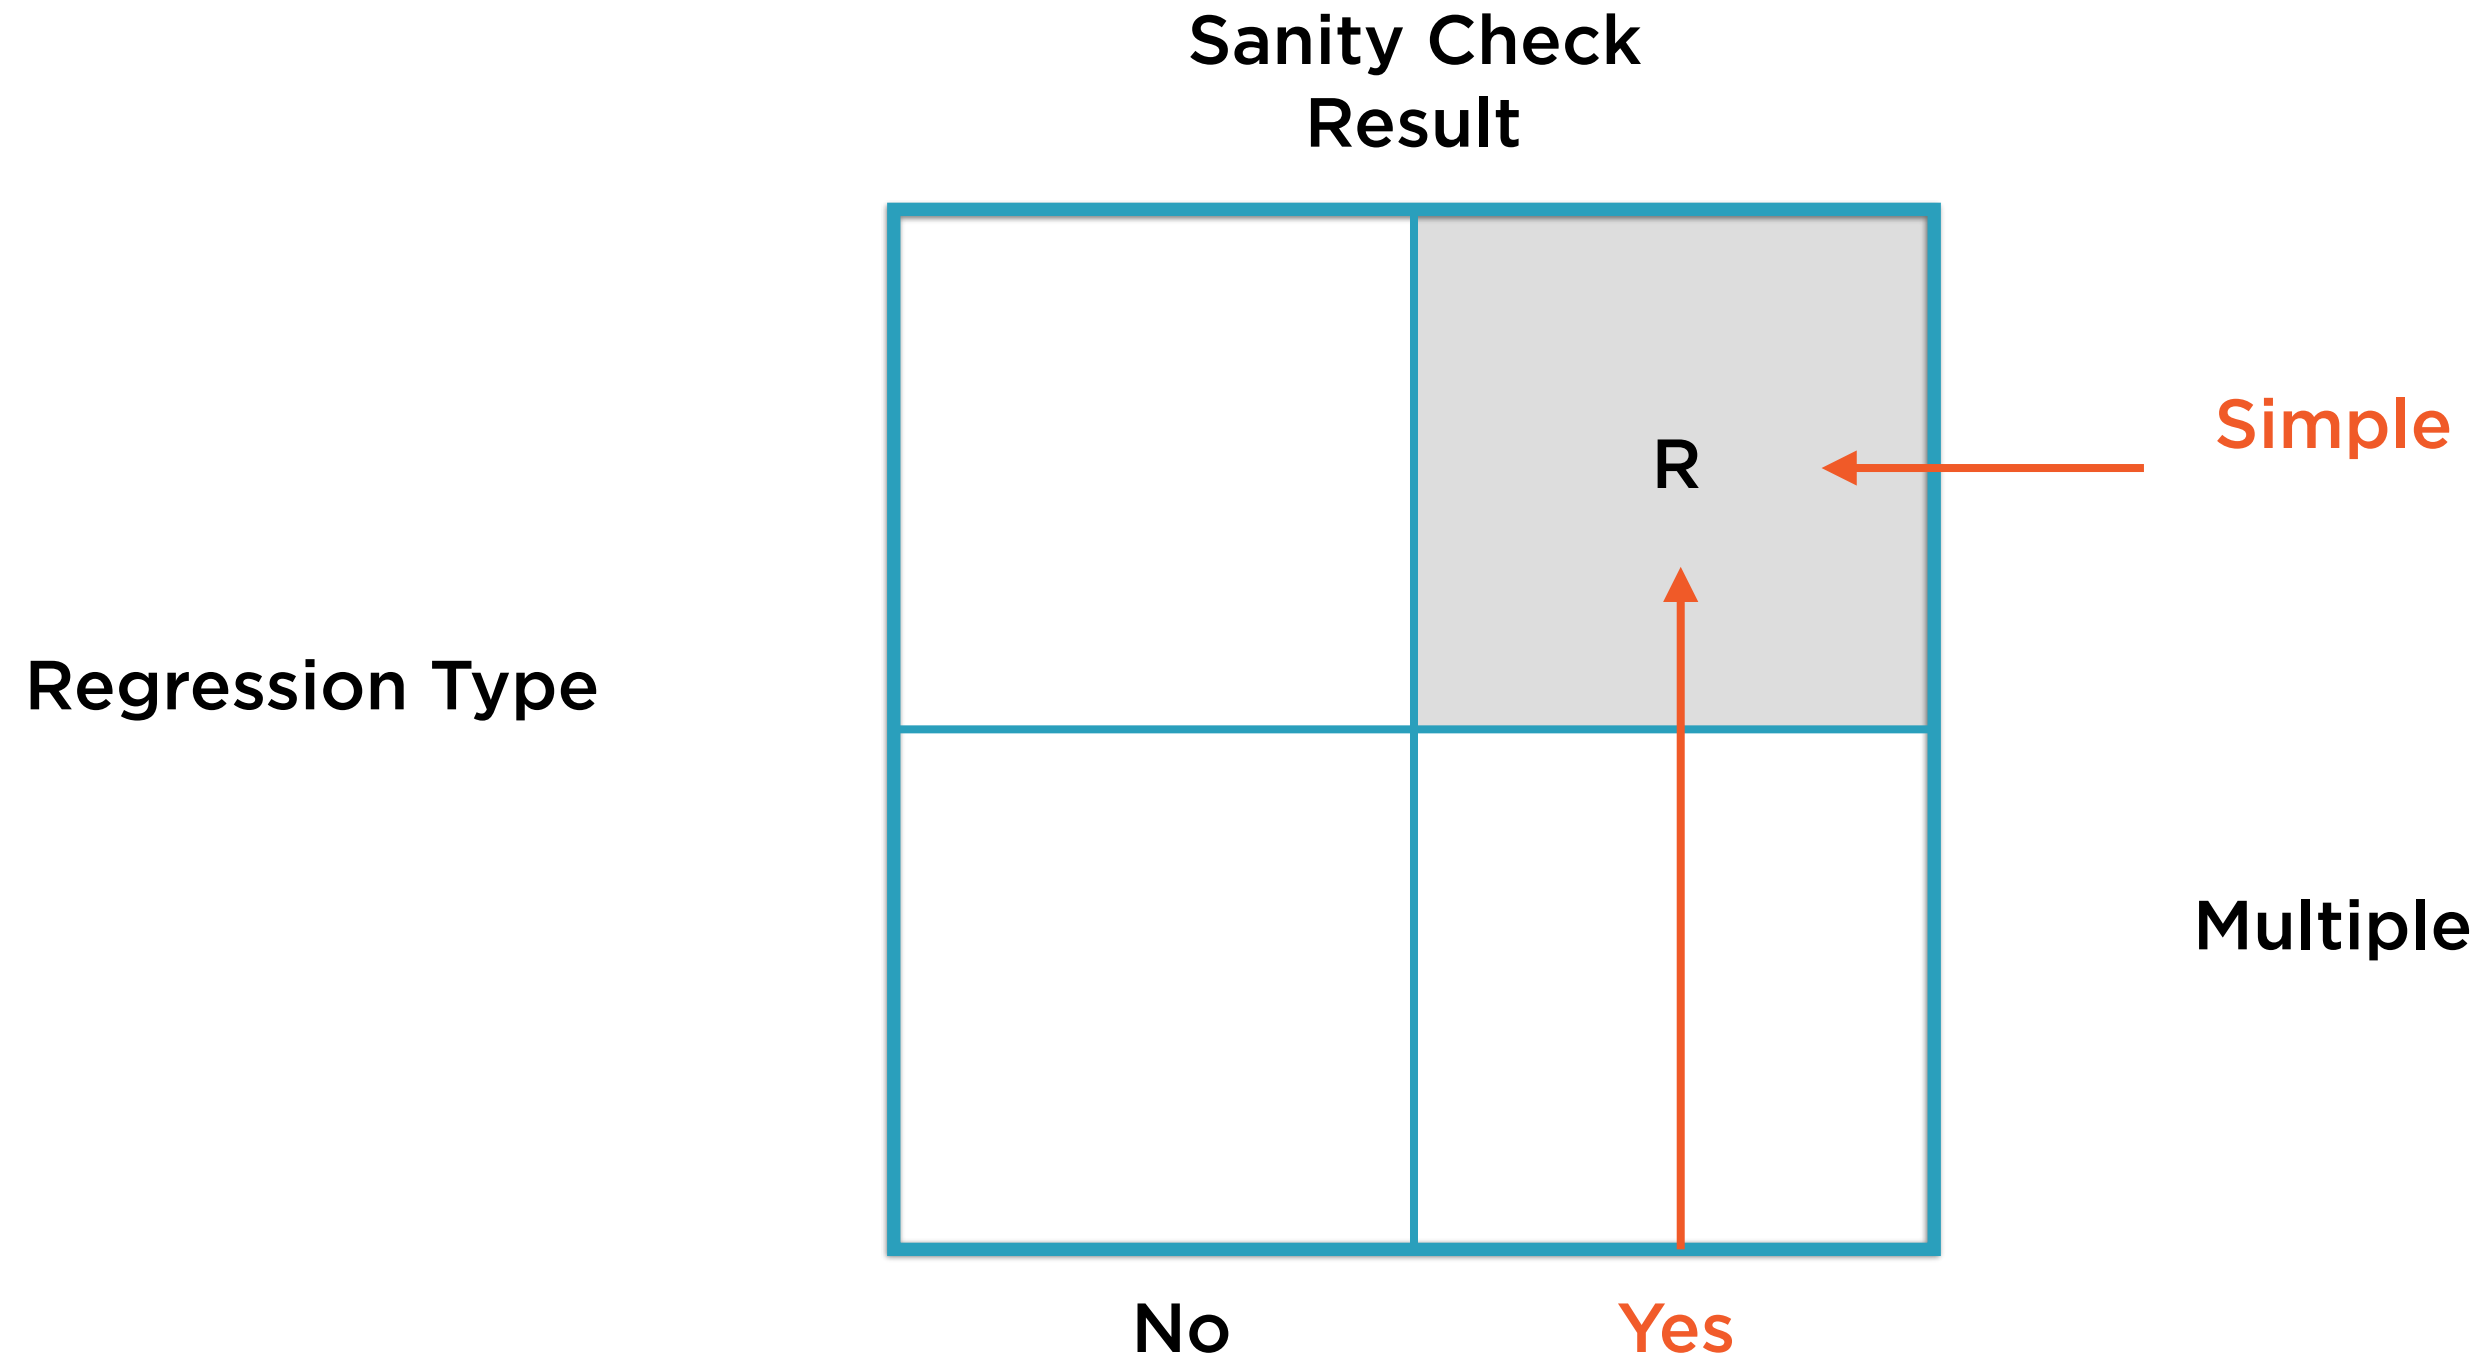
Excel or R for Linear Regression?




Excel or R for Linear Regression?



Excel or R for Linear Regression?



Excel or R for Linear Regression?



Excel or R for Linear Regression?

Sanity Check Result			
Regression Type	No	Yes	
	Excel	R	Simple
	R	R	Multiple

Demo

Simple regression models in R

Data Frame: Data in Rows and Columns

Each row represents 1 observation	DATE	OPEN	...	ADJUSTED CLOSE	Each column represents 1 variable (a list or vector)
	2016-12-01	772	...	779	
	2016-11-01	758	...	747	
	2006-01-01	302	...	309	

From File to Data Frame

DATE	OPEN	...	ADJUSTED CLOSE
2016-12-01	772	...	779
2016-11-01	758	...	747
2006-01-01	302	...	309

File

→
`read.table`

DATE	OPEN	...	ADJUSTED CLOSE
2016-12-01	772	...	779
2016-11-01	758	...	747
2006-01-01	302	...	309

Data Frame

Stripping Irrelevant Columns

DATE	OPEN	...	ADJUSTED CLOSE
2016-12-01	772	...	779
2016-11-01	758	...	747
2006-01-01	302	...	309

Data Frame 1

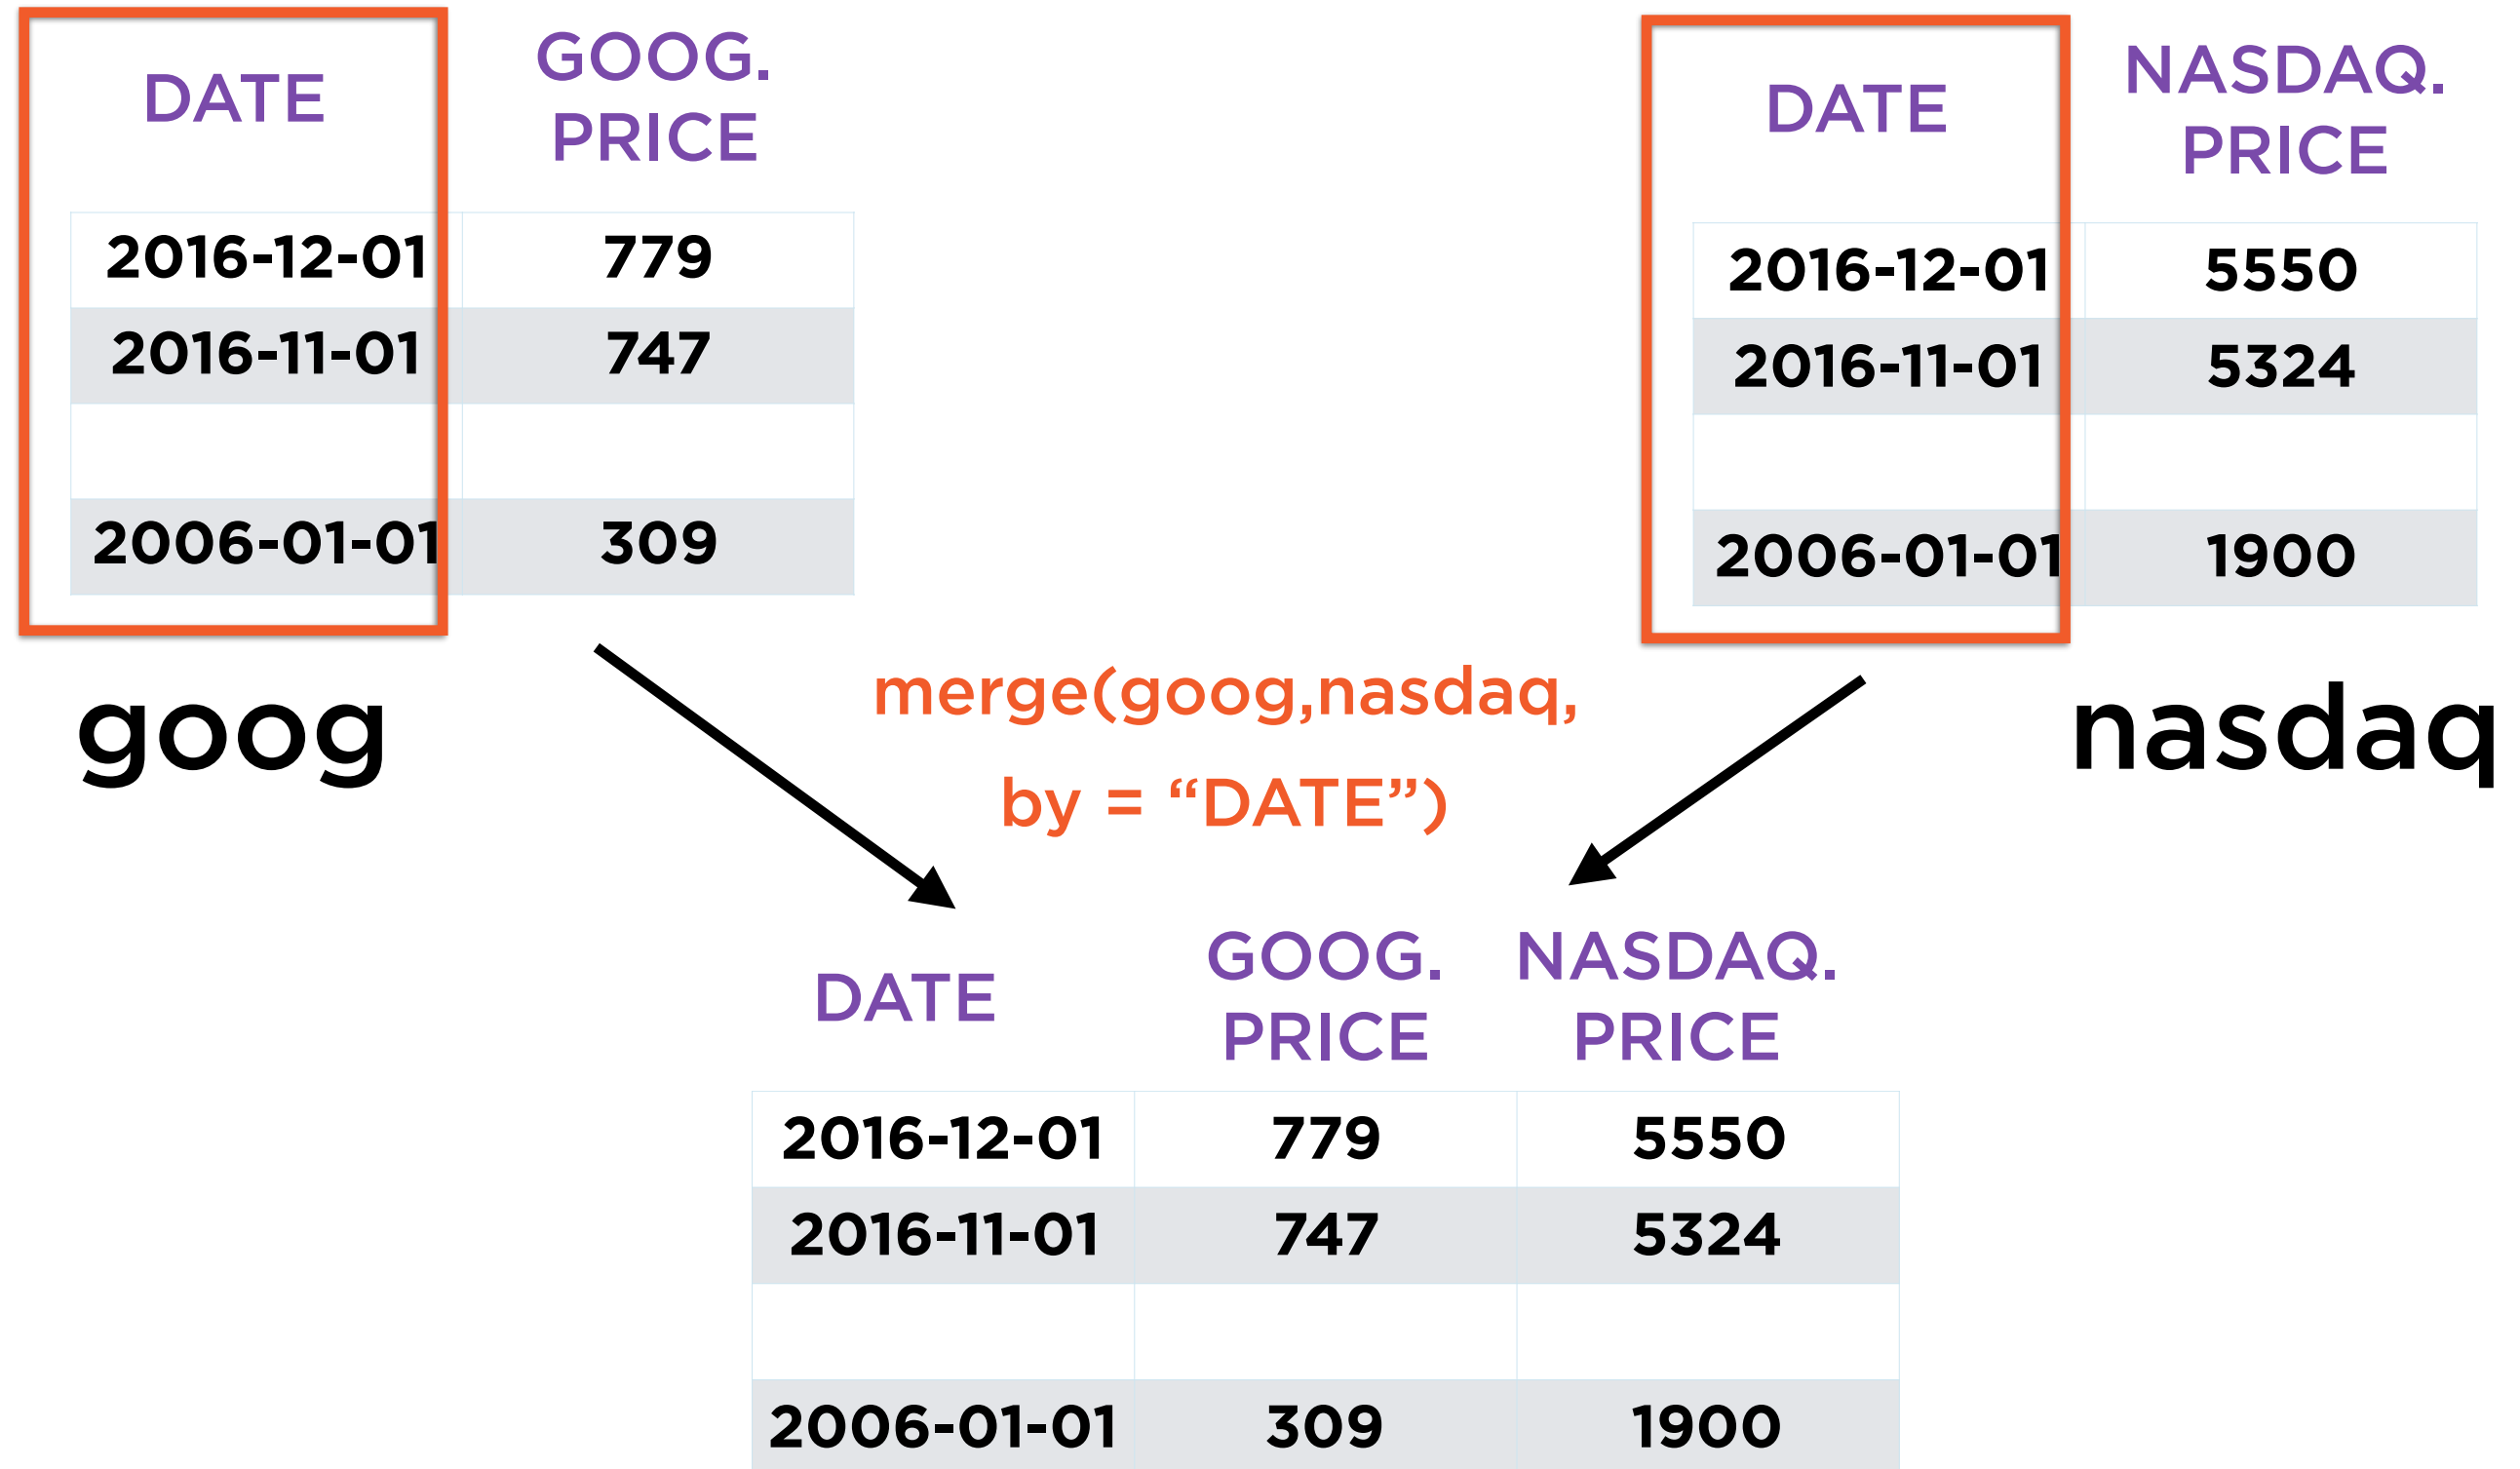


`[,c("DATE", "ADJUSTED
CLOSE")]`

DATE	ADJUSTED CLOSE
2016-12-01	779
2016-11-01	747
2006-01-01	309

Data Frame 2

Merging Data Frames



Negative Indices => Exclude Data

goog

DATE	GOOG. PRICE	NASDAQ. PRICE
2016-12-01	779	5550
2016-11-01	747	5324
2006-01-01	309	1900

Row 1

Row nrow(goog)

Column 1

goog[-nrow(goog),-1]

Negative Indices => Exclude Data

goog

DATE	GOOG. PRICE	NASDAQ. PRICE
2016-12-01	779	5550
2016-11-01	747	5324
2006-01-01	309	1900

Row 1

Row nrow(goog)

Column 1

goog[-nrow(goog),-1]

Negative Indices => Exclude Data

goog

DATE	GOOG. PRICE	NASDAQ. PRICE
2016-12-01	779	5550
2016-11-01	747	5324
2006-01-01	309	1900

Exclude

Column 1

Row 1

Row nrow(goog)

`goog[-nrow(goog),-1]`

Negative Indices => Exclude Data

goog

Exclude

DATE	GOOG. PRICE	NASDAQ. PRICE	
2016-12-01	779	5550	Row 1
2016-11-01	747	5324	
2006-01-01	309	1900	Row nrow(goog)

Column 1

`goog[-nrow(goog), -1]`

Negative Indices => Exclude Data

goog

	DATE	GOOG. PRICE	NASDAQ. PRICE	
	2016-12-01	779	5550	Row 1
	2016-11-01	747	5324	
Exclude	2006-01-01	309	1900	Row nrow(goog)

Column 1

goog[-nrow(goog),-1]

Element-wise Operations

779	5550	/	747	5324	=	779/747	5550/5324
					
					

**goog[-nrow(goog),-1]/
goog[-1,-1]**

Element-wise Operations

779/747	5550/5324		1	1		779/747 - 1	5550/5324 - 1
...	...		1	1	
		-	1	1	=		
			1	1			
...	...		1	1	

`goog[-nrow(goog),-1]/`
`goog[-1,-1] - 1`

This converts prices to returns

Summary

Built regression models in R

Avoided some common regression pitfalls

Use simple regression models in R

- **to explain variance**
- **to make forecasts**